

10_Flights_Challenge

September 4, 2021

Flights Challenge

1 Flights Data Exploration Challenge

```
[2]: import pandas as pd
```

```
df_flights = pd.read_csv('flights.csv')
df_flights.head()
```

```
[2]:   Year  Month  DayOfMonth  DayOfWeek  Carrier  OriginAirportID  \
0  2013     9         16           1      DL          15304
1  2013     9         23           1      WN          14122
2  2013     9          7           6      AS          14747
3  2013     7         22           1      OO          13930
4  2013     5         16           4      DL          13931
```

```
      OriginAirportName  OriginCity  OriginState  DestAirportID  \
0      Tampa International      Tampa          FL          12478
1  Pittsburgh International  Pittsburgh          PA          13232
2  Seattle/Tacoma International    Seattle          WA          11278
3  Chicago O'Hare International    Chicago          IL          11042
4      Norfolk International    Norfolk          VA          10397
```

```
      DestAirportName  DestCity  DestState  CRSDepTime  \
0  John F. Kennedy International  New York      NY          1539
1  Chicago Midway International    Chicago      IL           710
2  Ronald Reagan Washington National  Washington      DC           810
3  Cleveland-Hopkins International    Cleveland      OH           804
4  Hartsfield-Jackson Atlanta International    Atlanta      GA           545
```

```
      DepDelay  DepDel15  CRSArrTime  ArrDelay  ArrDel15  Cancelled
0           4         0.0        1824         13          0          0
1           3         0.0         740         22          1          0
2          -3         0.0        1614         -7          0          0
3          35         1.0        1027         33          1          0
4          -1         0.0         728         -9          0          0
```

```
[4]: len(df_flights.columns)
df_flights.shape
```

```
[4]: (271940, 20)
```

1.1 Finding Missing Data

1. Start by cleaning the data.

- Identify any null or missing data, and impute appropriate replacement values.
- Identify and eliminate any outliers in the **DepDelay** and **ArrDelay** columns.

```
[5]: missing_data = df_flights.isnull().sum()
missing_data
```

```
[5]: Year                0
Month                  0
DayofMonth             0
DayOfWeek              0
Carrier                0
OriginAirportID        0
OriginAirportName      0
OriginCity             0
OriginState            0
DestAirportID          0
DestAirportName        0
DestCity               0
DestState              0
CRSDepTime             0
DepDelay               0
DepDel15               2761
CRSArrTime             0
ArrDelay               0
ArrDel15               0
Cancelled              0
dtype: int64
```

2761 data points are missing from the DepDel15 feature which indicates if departure is delayed by more than 15 minutes

```
[7]: missing_data=missing_data.to_frame()
missing_data=missing_data.rename(columns={0: 'Empty Cells'})
print(missing_data)
```

```
Empty Cells
Year                0
Month              0
DayofMonth         0
DayOfWeek          0
Carrier            0
```

```

OriginAirportID      0
OriginAirportName     0
OriginCity            0
OriginState           0
DestAirportID         0
DestAirportName       0
DestCity              0
DestState             0
CRSDepTime            0
DepDelay              0
DepDel15              2761
CRSArrTime            0
ArrDelay              0
ArrDel15              0
Cancelled             0

```

1.2 checking for zeros

```
[18]: df_flights.isnull().any(axis=1)
```

```

[18]: 0      False
      1      False
      2      False
      3      False
      4      False
      ...
      271935  False
      271936  False
      271937  False
      271938  False
      271939  False
      Length: 271940, dtype: bool

```

since only DepDel15 has missing value, lets describe the subset of dataset that has missing value, thus we are picking 2761 rows with NaN seen in DepDel15

```
[21]: df_flights[df_flights.isnull().any(axis=1)].describe()
```

```

[21]:
   count  Year      Month  DayOfMonth  DayOfWeek  OriginAirportID  \
count  2761.0  2761.000000  2761.000000  2761.000000  2761.000000
mean   2013.0    6.455632   15.572619    3.604853   12757.763129
std      0.0    1.759942    8.092708    1.748487   1426.462196
min   2013.0    4.000000    1.000000    1.000000   10140.000000
25%   2013.0    5.000000   10.000000    2.000000   11298.000000
50%   2013.0    6.000000   16.000000    4.000000   12892.000000
75%   2013.0    8.000000   22.000000    5.000000   13930.000000
max   2013.0   10.000000   31.000000    7.000000   15376.000000

```

	DestAirportID	CRSDepTime	DepDelay	DepDel15	CRSArrTime	ArrDelay	\
count	2761.000000	2761.000000	2761.0	0.0	2761.000000	2761.0	
mean	12708.952553	1431.354944	0.0	NaN	1587.419051	0.0	
std	1408.166022	457.450773	0.0	NaN	485.236232	0.0	
min	10140.000000	5.000000	0.0	NaN	5.000000	0.0	
25%	11298.000000	1050.000000	0.0	NaN	1229.000000	0.0	
50%	12892.000000	1500.000000	0.0	NaN	1645.000000	0.0	
75%	13930.000000	1815.000000	0.0	NaN	2003.000000	0.0	
max	15376.000000	2359.000000	0.0	NaN	2359.000000	0.0	

	ArrDel15	Cancelled
count	2761.0	2761.0
mean	1.0	1.0
std	0.0	0.0
min	1.0	1.0
25%	1.0	1.0
50%	1.0	1.0
75%	1.0	1.0
max	1.0	1.0

Appart from DepDel15, notice DepDelay has 0 as min, max and mode, but DepDel15 is same but has NaN instead. > **DepDelay** is the number of minutes departure was delayed > > **DepDel15** feature which indicates if departure is delayed by more than 15 minutes > > Since Since there is no delay showing 0 in DepDelay, there is no 15 minutes delay. > > Replay all NaN in DepDel15 with 0.

```
[22]: df_flights.DepDel15 = df_flights.DepDel15.fillna(0)
df_flights.isnull().sum()
```

```
[22]: Year          0
      Month         0
      DayOfMonth    0
      DayOfWeek     0
      Carrier       0
      OriginAirportID 0
      OriginAirportName 0
      OriginCity     0
      OriginState    0
      DestAirportID  0
      DestAirportName 0
      DestCity       0
      DestState      0
      CRSDepTime     0
      DepDelay       0
      DepDel15       0
      CRSArrTime     0
      ArrDelay       0
      ArrDel15       0
```

Cancelled 0
dtype: int64

```
[23]: df_flights['DepDelay'].name # shows the name of the column
```

```
[23]: 'DepDelay'
```

- identify and eliminate any outliers in the **DepDelay** and **ArrDelay** columns.

1.3 Clean outliers

View the distribution and summary for the DepDelay and ArrDelay columns.

```
[24]: def show_distribution(var_data):  
    from matplotlib import pyplot as plt  
  
    # Get statistics  
    min_val = var_data.min()  
    max_val = var_data.max()  
    mean_val = var_data.mean()  
    med_val = var_data.median()  
    mod_val = var_data.mode()[0]  
  
    print(var_data.name, '\nMinimum:{:.2f}\nMean:{:.2f}\nMedian:{:.2f}\nMode:{:.2f}\nMaximum:{:.2f}\n'.format(min_val,  
→ mean_val,  
→ med_val,  
→ mod_val,  
→ max_val))  
  
    # Create a figure for 2 subplots (2 rows, 1 column)  
    fig, ax = plt.subplots(2, 1, figsize = (10, 4))  
  
    # Plot the histogram  
    ax[0].hist(var_data)  
    ax[0].set_ylabel('Frequency')  
  
    # Add lines for mean, median, and mode  
    ax[0].axvline(x=min_val, color = 'gray', linestyle='dashed', linewidth= 2)  
    ax[0].axvline(x=mean_val, color = 'cyan', linestyle='dashed', linewidth = 2)  
    ax[0].axvline(x=med_val, color = 'red', linestyle='dashed', linewidth = 2)  
    ax[0].axvline(x=mod_val, color = 'yellow', linestyle='dashed', linewidth = 2)  
→  
    ax[0].axvline(x=max_val, color = 'gray', linestyle='dashed', linewidth = 2)
```

```

# Plot the boxplot
ax[1].boxplot(var_data, vert=False)
ax[1].set_xlabel('Value')

# Add a title to the Figure
fig.suptitle(var_data.name)

# Show the figure
fig.show()

```

```

[25]: # Call the function for each delay field
delayFields = ['DepDelay', 'ArrDelay']
for col in delayFields:
    show_distribution(df_flights[col])

```

```

DepDelay
Minimum:-63.00
Mean:10.35
Median:-1.00
Mode:-3.00
Maximum:1425.00

```

```

C:\Users\aduzo\Anaconda3\lib\site-packages\ipykernel_launcher.py:39:
UserWarning: Matplotlib is currently using
module://ipykernel.pylab.backend_inline, which is a non-GUI backend, so cannot
show the figure.

```

```

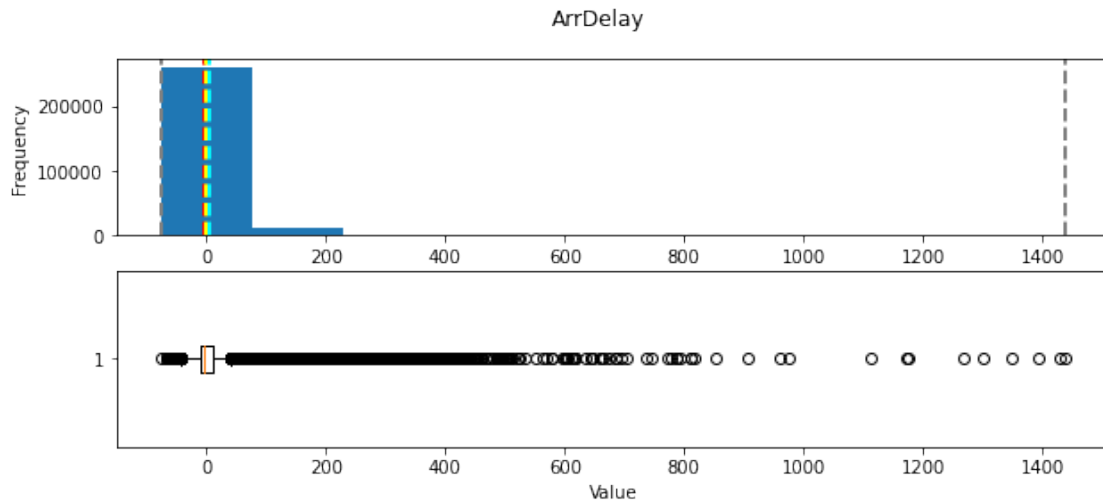

ArrDelay
Minimum:-75.00
Mean:6.50
Median:-3.00
Mode:0.00
Maximum:1440.00

```

```

C:\Users\aduzo\Anaconda3\lib\site-packages\ipykernel_launcher.py:39:
UserWarning: Matplotlib is currently using
module://ipykernel.pylab.backend_inline, which is a non-GUI backend, so cannot
show the figure.

```



There are a outliers at the lower and upper ends of both variables - particularly at the upper end.

Let's trim the data so that we include only rows where the values for these fields are within the 1st and 90th percentile.

```
[26]: # Trim outliers for ArrDelay based on 1% and 90% percentiles
ArrDelay_01pcntile = df_flights.ArrDelay.quantile(0.01)
ArrDelay_90pcntile = df_flights.ArrDelay.quantile(0.90)
print(f'1% percentile is {ArrDelay_01pcntile} \n90% percentile is {ArrDelay_90pcntile}')
```

1% percentile is -33.0

90% percentile is 38.0

-33 is at 1% while 38 is at 90% of the data

```
[27]: df_flights.shape
```

```
[27]: (271940, 20)
```

```
[28]: df_flights = df_flights[df_flights.ArrDelay < ArrDelay_90pcntile]
df_flights = df_flights[df_flights.ArrDelay > ArrDelay_01pcntile]
```

```
[29]: df_flights.shape
```

```
[29]: (241916, 20)
```

```
[30]: # Trim outliers for DepDelay based on 1% and 90% percentiles
DepDelay_01pcntile = df_flights.DepDelay.quantile(0.01)
DepDelay_90pcntile = df_flights.DepDelay.quantile(0.90)
df_flights = df_flights[df_flights.DepDelay < DepDelay_90pcntile]
df_flights = df_flights[df_flights.DepDelay > DepDelay_01pcntile]
```



```
[31]: # View the revised distributions
for col in delayFields:
    show_distribution(df_flights[col])
```

```
DepDelay
Minimum:-11.00
Mean:-0.92
Median:-2.00
Mode:-3.00
Maximum:16.00
```

```
C:\Users\aduzo\Anaconda3\lib\site-packages\ipykernel_launcher.py:39:
UserWarning: Matplotlib is currently using
module://ipykernel.pylab.backend_inline, which is a non-GUI backend, so cannot
show the figure.
```

```
ArrDelay
Minimum:-32.00
Mean:-5.03
Median:-6.00
Mode:0.00
Maximum:37.00
```

```
C:\Users\aduzo\Anaconda3\lib\site-packages\ipykernel_launcher.py:39:
UserWarning: Matplotlib is currently using
module://ipykernel.pylab.backend_inline, which is a non-GUI backend, so cannot
show the figure.
```

```
[32]: df_flights.shape
```

```
[32]: (214397, 20)
```

That looks a bit better

2. Explore the cleaned data.

- View summary statistics for the numeric fields in the dataset.
- Determine the distribution of the **DepDelay** and **ArrDelay** columns.
- Use statistics, aggregate functions, and visualizations to answer the following questions:
 - *What are the average (mean) departure and arrival delays?*
 - *How do the carriers compare in terms of arrival delay performance?*
 - *Are some days of the week more prone to arrival days than others?*

- Which departure airport has the highest average departure delay?
- Do **late** departures tend to result in longer arrival delays than on-time departures?
- Which route (from origin airport to destination airport) has the most **late** arrivals?
- Which route has the highest average arrival delay?

1.3.1 Explore the data

Let's start with an overall view of the summary statistics for the numeric columns.

```
[33]: df_flights.describe()
```

```
[33]:
```

	Year	Month	DayofMonth	DayOfWeek	OriginAirportID	\
count	214397.0	214397.000000	214397.000000	214397.000000	214397.000000	
mean	2013.0	7.018368	15.794703	3.902737	12757.827661	
std	0.0	2.006398	8.859118	1.997744	1510.058629	
min	2013.0	4.000000	1.000000	1.000000	10140.000000	
25%	2013.0	5.000000	8.000000	2.000000	11292.000000	
50%	2013.0	7.000000	16.000000	4.000000	12892.000000	
75%	2013.0	9.000000	23.000000	6.000000	14100.000000	
max	2013.0	10.000000	31.000000	7.000000	15376.000000	

	DestAirportID	CRSDepTime	DepDelay	DepDel15	\
count	214397.000000	214397.000000	214397.000000	214397.000000	
mean	12726.276147	1278.223879	-0.921692	0.018116	
std	1506.251757	469.440262	5.708594	0.133371	
min	10140.000000	1.000000	-11.000000	0.000000	
25%	11292.000000	850.000000	-5.000000	0.000000	
50%	12892.000000	1235.000000	-2.000000	0.000000	
75%	14057.000000	1655.000000	1.000000	0.000000	
max	15376.000000	2359.000000	16.000000	1.000000	

	CRSArrTime	ArrDelay	ArrDel15	Cancelled
count	214397.000000	214397.000000	214397.000000	214397.000000
mean	1461.406596	-5.030276	0.068602	0.013228
std	485.676457	11.424280	0.252776	0.114249
min	1.000000	-32.000000	0.000000	0.000000
25%	1054.000000	-13.000000	0.000000	0.000000
50%	1445.000000	-6.000000	0.000000	0.000000
75%	1845.000000	2.000000	0.000000	0.000000
max	2359.000000	37.000000	1.000000	1.000000

```
[34]: delayFields
```

```
[34]: ['DepDelay', 'ArrDelay']
```

What are the mean departure and arrival delays?

```
[35]: df_flights[delayFields].mean()
```

```
[35]: DepDelay    -0.921692  
      ArrDelay    -5.030276  
      dtype: float64
```

How do the carriers compare in terms of arrival delay performance?


```
[36]: df_flights.Carrier.unique()
```


```
[36]: array(['DL', 'WN', 'AS', 'EV', 'AA', 'OO', 'US', 'UA', 'FL', 'B6', 'VX',  
          'MQ', '9E', 'YV', 'F9', 'HA'], dtype=object)
```

```
[39]: print(f'{len(df_flights.Carrier.unique())} unique Carriers')
```

```
[39]: 16
```

```
[40]: for col in delayFields:  
      df_flights.boxplot(column=col, by='Carrier', figsize=(8,8))
```





As seen above, for DepDelay most carriers have a mean with no delay or depart earlier on average, thus no delay, but delays occur, but in some instances, departure was delayed by more than 5 minutes. Carriers such as 9E, AS, HA 75% of their times depart earlier than the scheduled time.

For ArrDelay, all carriers arrive on time on average.

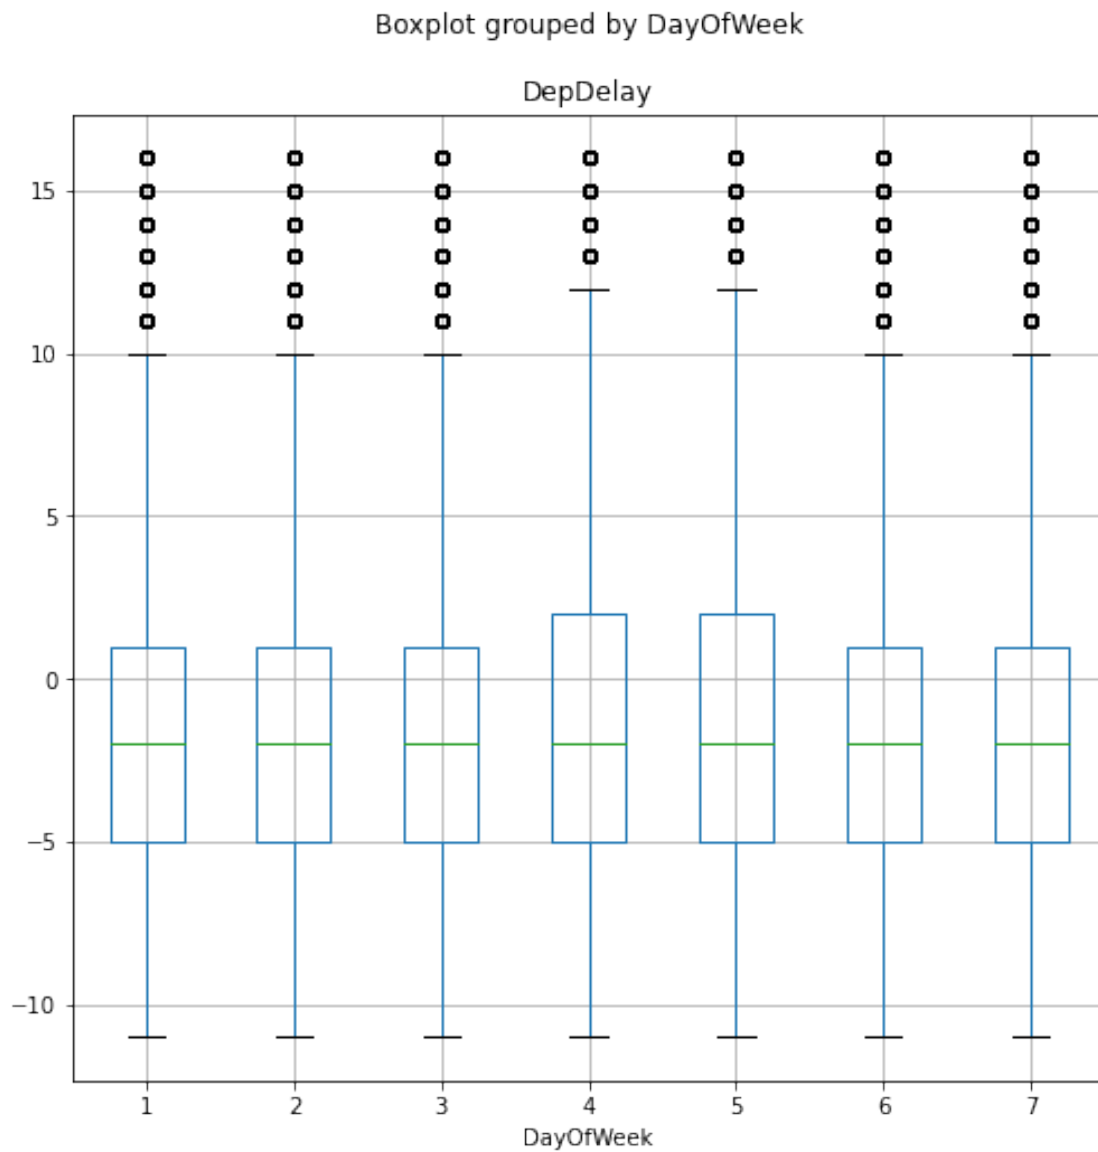
Are some days of the week more prone to arrival delays than others?


```
[42]: df_flights.DayOfWeek.unique()
```

```
[42]: array([1, 6, 4, 7, 2, 5, 3], dtype=int64)
```

DayOfWeek: The day of the week on which the flight departed - from 1 (Monday) to 7 (Sunday)

```
[41]: for col in delayFields:
      df_flights.boxplot(column=col, by='DayOfWeek', figsize=(8,8))
```





All days of the week has on average same arrival time. But Thursday is slightly more prone, but this is negligible.

Which departure airport has the highest average departure delay?

```
[44]: print(f'There are {len(df_flights.OriginAirportName.unique())} departure_  
      ↪airports')
```

There are 70 departure airports


```
[45]: departure_airport_group = df_flights.groupby(df_flights.OriginAirportName)
```

```
mean_departure_delays =pd.DataFrame(departure_airport_group['DepDelay'].mean()).
↳sort_values('DepDelay', ascending=False)
mean_departure_delays.plot(kind = 'bar', figsize=(12,12))
mean_departure_delays
```

```
[45]:
```

OriginAirportName	DepDelay
Chicago Midway International	2.365960
Dallas Love Field	2.148798
William P Hobby	1.561927
Metropolitan Oakland International	0.964853
Denver International	0.807272
...	...
Memphis International	-2.962737
Tucson International	-2.989154
Ted Stevens Anchorage International	-3.149758
Richmond International	-3.198073
Long Beach Airport	-3.447844

[70 rows x 1 columns]



Do late departures tend to result in longer arrival delays than on-time departures?


```
[47]: df_flights.DepDel15.unique()
```

```
[47]: array([0., 1.])
```

0 if no delay by 15 minutes or more, 1 if so

```
[50]: df_flights.boxplot(column='ArrDelay', by='DepDel15', figsize=(12,12))
```

```
[50]: <AxesSubplot:title={'center':'ArrDelay'}, xlabel='DepDel15'>
```



Yes as seen in the above box plot,

```
[51]: df_flights.groupby(by='DepDel15').ArrDelay.mean()
```

```
[51]: DepDel15
      0.0    -5.299464
      1.0     9.559732
      Name: ArrDelay, dtype: float64
```

As seen from the above graph, late departure influences arrival time as shown in there mean, early departure arrives 5 minutes earlier on average, while late departure arrives 10 minutes later than scheduled.

Which route (from origin airport to destination airport) has the most late arrivals?

Note: no column named route, rather there is Origin and Destination columns, so route is by combining these two columns.

```
[52]: df_flights.columns
```

```
[52]: Index(['Year', 'Month', 'DayofMonth', 'DayOfWeek', 'Carrier',
        'OriginAirportID', 'OriginAirportName', 'OriginCity', 'OriginState',
        'DestAirportID', 'DestAirportName', 'DestCity', 'DestState',
        'CRSDepTime', 'DepDelay', 'DepDel15', 'CRSArrTime', 'ArrDelay',
        'ArrDel15', 'Cancelled'],
        dtype='object')
```

```
[53]: # Add a routes column

routes = pd.Series(df_flights['OriginAirportName']+
    ↪ '-'->'+df_flights['DestAirportName'])
df_flights = pd.concat([df_flights, routes.rename("Route")], axis=1)

# Group by routes
route_group = df_flights.groupby(df_flights.Route)
pd.DataFrame(route_group['ArrDel15'].sum()).
    ↪ sort_values('ArrDel15',ascending=False)
```

```
[53]:
```

Route	ArrDel15
San Francisco International->Los Angeles Intern...	90
Los Angeles International->San Francisco Intern...	69
LaGuardia->Hartsfield-Jackson Atlanta Internati...	68
Los Angeles International->John F. Kennedy Inte...	52
LaGuardia->Charlotte Douglas International	51
...	...
Logan International->Austin - Bergstrom Interna...	0
Logan International->Memphis International	0
Logan International->Port Columbus International	0
San Diego International->Cincinnati/Northern Ke...	0
Louis Armstrong New Orleans International->San ...	0

```
[2479 rows x 1 columns]
```

Which route has the highest average arrival delay?

```
[54]: pd.DataFrame(route_group['ArrDelay'].mean()).sort_values('ArrDelay',  
↪ascending=False)
```

```
[54]:
```

Route	ArrDelay
Louis Armstrong New Orleans International->Rona...	24.500000
Cleveland-Hopkins International->Palm Beach Int...	18.000000
John F. Kennedy International->Louisville Inter...	18.000000
Cleveland-Hopkins International->Philadelphia I...	12.800000
Memphis International->Denver International	9.758621
...	...
Lambert-St. Louis International->Cleveland-Hopk...	-20.000000
Eppley Airfield->LaGuardia	-20.750000
Denver International->Kahului Airport	-22.666667
Jacksonville International->Chicago Midway Inte...	-24.125000
Indianapolis International->Logan International	-26.000000

[2479 rows x 1 columns]